





h (token embeddings)



h



Inference – hard gating



no (active)


yes (halted)



Training – soft gating



$\Delta^{\text{attn}}, \Delta^{\text{ffn}}$



h (updated)